

EXAMTOPICS

- Expert Verified, Online, Free.

≡ MENU



🔍 Google Discussions



Exam Professional Machine Learning Engineer All Questions

View all questions & answers for the Professional Machine Learning Engineer exam

Go to Exam

EXAM PROFESSIONAL MACHINE LEARNING ENGINEER TOPIC 1 QUESTION 64 DISCUSSIO..

Actual exam question from Google's Professional Machine Learning Engineer

Question #: 64


Topic #: 1

[\[All Professional Machine Learning Engineer Questions\]](#)

You recently designed and built a custom neural network that uses critical dependencies specific to your organization's framework. You need to train the model using a managed training service on Google Cloud. However, the ML framework and related dependencies are not supported by AI Platform Training. Also, both your model and your data are too large to fit in memory on a single machine. Your ML framework of choice uses the scheduler, workers, and servers distribution structure. What should you do?

- A. Use a built-in model available on AI Platform Training.
- B. Build your custom container to run jobs on AI Platform Training.
- C. Build your custom containers to run distributed training jobs on AI Platform Training.
- D. Reconfigure your code to a ML framework with dependencies that are supported by AI Platform Training.

Show Suggested Answer

by  Vedjha at Dec. 7, 2022, 11:46 p.m.

Comments

Type your comment...

Submit

mil_spyro Highly Voted 1 year, 10 months ago

Selected Answer: C

Answer C. By running your machine learning (ML) training job in a custom container, you can use ML frameworks, non-ML dependencies, libraries, and binaries that are not otherwise supported on Vertex AI.

Model and your data are too large to fit in memory on a single machine hence distributed training jobs.

<https://cloud.google.com/vertex-ai/docs/training/containers-overview>

upvoted 9 times

PhilipKoku Most Recent 5 months ago

Selected Answer: C

C) Distributed training with customer containers

upvoted 1 times

MultiCloudIronMan 7 months, 1 week ago

Selected Answer: C

This allows using external dependences and distributed training will solve the memory issues

upvoted 3 times

Werner123 8 months, 1 week ago

Selected Answer: C

Critical dependencies that are not supported -> Custom container

Too large to fit in memory on a single machine -> Distributed

upvoted 2 times

M25 1 year, 6 months ago

Selected Answer: C

Went with C

upvoted 1 times

wish0035 1 year, 10 months ago

Selected Answer: C

ans: C

A, D => too much work.

B => discarded because "model and your data are too large to fit in memory on a single machine"

upvoted 1 times

ares81 1 year, 11 months ago

C, for me!

upvoted 1 times

JeanEl 1 year, 11 months ago

Selected Answer: C

I think it's C

upvoted 1 times

Vedjha 1 year, 11 months ago

Will go for 'C'- Custom containers can address the env limitation and distributed processing will handle the data volume

upvoted 1 times



Social Media

[Facebook](#) , [Twitter](#)

[YouTube](#) , [Reddit](#)

[Pinterest](#)



We are the biggest and most updated IT certification exam material website.

Using our own resources, we strive to strengthen the IT professionals community for free.



© 2024 ExamTopics

ExamTopics doesn't offer Real Microsoft Exam Questions. ExamTopics doesn't offer Real Amazon Exam Questions. ExamTopics Materials do not contain actual questions and answers from Cisco's Certification Exams.

CFA Institute does not endorse, promote or warrant the accuracy or quality of ExamTopics. CFA® and Chartered Financial Analyst® are registered trademarks owned by CFA Institute.