



- Expert Verified, Online, Free.

MENU



Google Discussions



## Exam Professional Machine Learning Engineer All Questions

View all questions & answers for the Professional Machine Learning Engineer exam

Go to Exam

### EXAM PROFESSIONAL MACHINE LEARNING ENGINEER TOPIC 1 QUESTION 105 DISCUSSI...

Actual exam question from Google's Professional Machine Learning Engineer

Question #: 105

Topic #: 1

[\[All Professional Machine Learning Engineer Questions\]](#)

You work for a gaming company that develops massively multiplayer online (MMO) games. You built a TensorFlow model that predicts whether players will make in-app purchases of more than \$10 in the next two weeks. The model's predictions will be used to adapt each user's game experience. User data is stored in BigQuery. How should you serve your model while optimizing cost, user experience, and ease of management?

- A. Import the model into BigQuery ML. Make predictions using batch reading data from BigQuery, and push the data to Cloud SQL
- B. Deploy the model to Vertex AI Prediction. Make predictions using batch reading data from Cloud Bigtable, and push the data to Cloud SQL.
- C. Embed the model in the mobile application. Make predictions after every in-app purchase event is published in Pub/Sub, and push the data to Cloud SQL.
- D. Embed the model in the streaming Dataflow pipeline. Make predictions after every in-app purchase event is published in Pub/Sub, and push the data to Cloud SQL.

Show Suggested Answer

by [hiromi](#) at Dec. 20, 2022, 9:57 p.m.

### Comments

Type your comment...

Submit


  **hiromi** Highly Voted  1 year, 10 months ago

**Selected Answer: A**

it seems A (not sure)

- <https://cloud.google.com/bigquery-ml/docs/reference/standard-sql/bigqueryml-syntax-create-tensorflow>

   upvoted 12 times

  **pinimichele01** Most Recent  6 months, 3 weeks ago

**Selected Answer: A**

Make predictions after every in-app purchase it is not necessary -> A

   upvoted 2 times

  **Mickey321** 11 months, 3 weeks ago

**Selected Answer: D**

Embedding the model in a streaming Dataflow pipeline allows low latency predictions on real-time events published to Pub/Sub. This provides a responsive user experience.

Dataflow provides a managed service to scale predictions and integrate with Pub/Sub, without having to manage servers.

Streaming predictions only when events occur optimizes cost compared to bulk or client-side prediction.

Pushing results to Cloud SQL provides a managed database for persistence.

In contrast, options A and B use inefficient batch predictions. Option C increases mobile app size and cost.


   upvoted 1 times

  **SamuelTsch** 1 year, 4 months ago

**Selected Answer: D**

D could be correct

   upvoted 1 times

  **Nxtgen** 1 year, 4 months ago

**Selected Answer: D**

These were my reasonings to choose D as best option:

B -> Vertex AI would not minimize cost

C -> Would not optimize user experience (this may lead to slow running of the game (lag)?)

A -> Would not optimize ease of management / automatization

D -> Best choice?

   upvoted 1 times

  **tavva\_prudhvi** 12 months ago

Why do you want to make a prediction after every app purchase bro?

   upvoted 3 times

  **M25** 1 year, 6 months ago

**Selected Answer: D**

For "used to adapt each user's game experience" points out to non-batch, hence excludes A & B, and embedding the model in the mobile app would not necessarily "optimize cost". Plus, the classical streaming solution builds on Dataflow along with Pub/Sub and BigQuery, embedding ML in Dataflow is low-code <https://cloud.google.com/blog/products/data-analytics/latest-dataflow-innovations-for-real-time-streaming-and-aiml> and apparently a modified version of the question points to the same direction <https://mikaelahonen.com/en/data/gcp-mle-exam-questions/>

   upvoted 3 times

  **ciro\_li** 1 year, 3 months ago

there's no need to make a prediction after every in-app purchase event. Am i wrong?

   upvoted 3 times

  **TNT87** 1 year, 6 months ago

**Selected Answer: A**

Yeah its A

   upvoted 2 times

  **TNT87** 1 year, 8 months ago

**Selected Answer: C**

Answer C

   upvoted 2 times

 **tavva\_prudhvi** 1 year, 7 months ago

Option C, embedding the model in the mobile application, can increase the size of the application and may not be suitable for real-time prediction.

   upvoted 2 times

[Start Learning for free](#)



### Social Media

[Facebook](#) , [Twitter](#)

[YouTube](#) , [Reddit](#)

[Pinterest](#)



We are the biggest and most updated IT certification exam material website.

Using our own resources, we strive to strengthen the IT professionals community for free.



© 2024 ExamTopics

ExamTopics doesn't offer Real Microsoft Exam Questions. ExamTopics doesn't offer Real Amazon Exam Questions. ExamTopics Materials do not

contain actual questions and answers from Cisco's Certification Exams.

CFA Institute does not endorse, promote or warrant the accuracy or quality of ExamTopics. CFA® and Chartered Financial Analyst® are registered trademarks owned by CFA Institute.