

EXAMTOPICS

- Expert Verified, Online, Free.

≡ MENU



🔍 Google Discussions



Exam Professional Machine Learning Engineer All Questions

View all questions & answers for the Professional Machine Learning Engineer exam

Go to Exam

EXAM PROFESSIONAL MACHINE LEARNING ENGINEER TOPIC 1 QUESTION 118 DISCUSSI...

Actual exam question from Google's Professional Machine Learning Engineer

Question #: 118

Topic #: 1

[\[All Professional Machine Learning Engineer Questions\]](#)

You are building a TensorFlow model for a financial institution that predicts the impact of consumer spending on inflation globally. Due to the size and nature of the data, your model is long-running across all types of hardware, and you have built frequent checkpointing into the training process. Your organization has asked you to minimize cost. What hardware should you choose?

- A. A Vertex AI Workbench user-managed notebooks instance running on an n1-standard-16 with 4 NVIDIA P100 GPUs
- B. A Vertex AI Workbench user-managed notebooks instance running on an n1-standard-16 with an NVIDIA P100 GPU
- C. A Vertex AI Workbench user-managed notebooks instance running on an n1-standard-16 with a non-preemptible v3-8 TPU
- D. A Vertex AI Workbench user-managed notebooks instance running on an n1-standard-16 with a preemptible v3-8 TPU

Show Suggested Answer

by  mymy9418 at Dec. 18, 2022, 2:47 a.m.

Comments

Type your comment...

Submit

🗒️ 👤 **hiromi** Highly Voted 1 year, 10 months ago

Selected Answer: D

D

you have built frequent checkpointing into the training process / minimize cost -> preemptible

👍 ↩️ 🚩 upvoted 7 times

🗒️ 👤 **5a74493** Most Recent 1 month, 2 weeks ago

Selected Answer: C

For financial institutions, reliability and minimizing interruptions are crucial. While preemptible instances are cost-effective, they do come with the risk of being terminated unexpectedly, which might not be ideal for critical financial applications.

👍 ↩️ 🚩 upvoted 1 times

🗒️ 👤 **M25** 1 year, 5 months ago

Selected Answer: D

Follows same principle as #70

👍 ↩️ 🚩 upvoted 2 times

🗒️ 👤 **Antmal** 1 year, 6 months ago

Selected Answer: D

Preemptible v3-8 TPUs are the most cost-effective option for training large TensorFlow models. They are up to 80% cheaper than non-preemptible v3-8 TPUs, and they are only preempted if Google Cloud needs the resources for other workloads.

In this case, the model is long-running and checkpointing is used. This means that the training process can be interrupted and resumed without losing any progress. Therefore, preemptible TPUs are a safe choice, as the training process will not be interrupted if the TPU is preempted.

The other options are not as cost-effective.

👍 ↩️ 🚩 upvoted 1 times

🗒️ 👤 **TNT87** 1 year, 7 months ago

Selected Answer: D

Answer D

👍 ↩️ 🚩 upvoted 1 times

🗒️ 👤 **ares81** 1 year, 9 months ago

Selected Answer: D

Frequent checkpoints --> Preemptible --> D

👍 ↩️ 🚩 upvoted 4 times

🗒️ 👤 **mymy9418** 1 year, 10 months ago

Selected Answer: D

preemptible is the keyword to me

👍 ↩️ 🚩 upvoted 1 times

Start Learning for free



Social Media

[Facebook](#) , [Twitter](#)

[YouTube](#) , [Reddit](#)

[Pinterest](#)



We are the biggest and most updated IT certification exam material website.

Using our own resources, we strive to strengthen the IT professionals community for free.



© 2024 ExamTopics

ExamTopics doesn't offer Real Microsoft Exam Questions. ExamTopics doesn't offer Real Amazon Exam Questions. ExamTopics Materials do not contain actual questions and answers from Cisco's Certification Exams.

CFA Institute does not endorse, promote or warrant the accuracy or quality of ExamTopics. CFA® and Chartered Financial Analyst® are registered trademarks owned by CFA Institute.