

Google Discussions



Exam Professional Machine Learning Engineer All Questions

View all questions & answers for the Professional Machine Learning Engineer exam

Go to Exam

EXAM PROFESSIONAL MACHINE LEARNING ENGINEER TOPIC 1 QUESTION 173 DISCUSSI...

Actual exam question from Google's Professional Machine Learning Engineer

Question #: 173

Topic #: 1

[\[All Professional Machine Learning Engineer Questions\]](#)

You received a training-serving skew alert from a Vertex AI Model Monitoring job running in production. You retrained the model with more recent training data, and deployed it back to the Vertex AI endpoint, but you are still receiving the same alert. What should you do?

- A. Update the model monitoring job to use a lower sampling rate.
- B. Update the model monitoring job to use the more recent training data that was used to retrain the model.
- C. Temporarily disable the alert. Enable the alert again after a sufficient amount of new production traffic has passed through the Vertex AI endpoint.
- D. Temporarily disable the alert until the model can be retrained again on newer training data. Retrain the model again after a sufficient amount of new production traffic has passed through the Vertex AI endpoint.

Show Suggested Answer

by [b1a8fae](#) at Jan. 8, 2024, 5:17 p.m.

Comments

Type your comment...

Submit

  **AzureDP900** 4 months ago

C. Temporarily disable the alert. Enable the alert again after a sufficient amount of new production traffic has passed through the Vertex AI endpoint.

Here's why:

You've already retrained the model with more recent training data and deployed it back to the Vertex AI endpoint, but the alert persists.

This suggests that the model is still adapting to the changing data distribution in production.

Temporarily disabling the alert will give the model a chance to adjust to the new data distribution before the monitoring job starts firing alerts again.

Once enough new traffic has passed through, you can re-enable the alert and continue monitoring the model's performance.

   upvoted 1 times

  **info_appsatori** 4 months, 1 week ago

Selected Answer: B

The baseline is calculated when you create a Vertex AI Model Monitoring job, and is only recalculated if you update the training dataset for the job.

   upvoted 1 times

  **SahandJ** 5 months, 2 weeks ago

Is B actually the correct answer? According to the documentation, training-serving skew detection can only be enabled if the original training data is available. Furthermore, the baseline is automatically recalculated when the training data is updated.

So does this question imply that the model is trained on data without updating the original training-dataset? If so then B is clearly correct. If they updated the training dataset with new data and then retrained the model then the model monitoring job's baseline should automatically have been recalculated. I see no other valid answers in that case?

   upvoted 2 times

  **pinimichele01** 6 months, 1 week ago

Selected Answer: B

This option can help align the baseline distribution of the model monitoring job with the current distribution of the production data, and eliminate the false positive alerts.

   upvoted 1 times

  **36bdc1e** 9 months, 1 week ago

B

This option can help align the baseline distribution of the model monitoring job with the current distribution of the production data, and eliminate the false positive alerts.

   upvoted 3 times

  **BlehMaks** 9 months, 2 weeks ago

Selected Answer: B

the cause of the issue could be that the developer forgot to switch their monitoring job to the latest training dataset and the monitoring job still compares prod data with old training dataset and they of course have a skew

   upvoted 3 times

  **pikachu007** 9 months, 2 weeks ago

Selected Answer: B

B. Update the model monitoring job to use the more recent training data that was used to retrain the model:

This option directly aligns the model monitoring with the recently retrained model and ensures that the monitoring job reflects the characteristics of the latest training data.

   upvoted 1 times

  **b1a8fae** 9 months, 2 weeks ago

Selected Answer: D

A. Changing the sampling rate affects not training skew but cost efficiency: <https://cloud.google.com/vertex-ai/docs/model-monitoring/overview#considerations>

B. The model monitoring job is already using the most recent data to detect skew.

C&D are the same, except for D being more specific, so I would tend towards D.

   upvoted 1 times

Start Learning for free



Social Media

[Facebook](#) , [Twitter](#)

[YouTube](#) , [Reddit](#)

[Pinterest](#)



We are the biggest and most updated IT certification exam material website.

Using our own resources, we strive to strengthen the IT professionals community for free.



© 2024 ExamTopics

ExamTopics doesn't offer Real Microsoft Exam Questions. ExamTopics doesn't offer Real Amazon Exam Questions. ExamTopics Materials do not contain actual questions and answers from Cisco's Certification Exams.

CFA Institute does not endorse, promote or warrant the accuracy or quality of ExamTopics. CFA® and Chartered Financial Analyst® are registered trademarks owned by CFA Institute.